

Reproducible Research - Assignment 2 - NOAA storm database analysis

Synopsis

This document looks at the National Oceanic and Atmospheric Administration's (NOAA) Storm Data, collected from 1950 to 2011. The analysis calculates the storm event types that cause the most human damage (as measured by fatalities and injuries) and the storm event types that cause the most economic damage (as measured by crop and property damage).

Human damage: 'TORNADO' are by far the leading cause of death and injury, followed by 'HEAT' with less than 1/8th the fatalities and injuries, then 'THUNDERSTORM' and 'FLOOD'

Economic damage: 'FLOOD' causes by far the largest amount of economic damage, followed by 'HURRICANE' with about half as much, then 'TORNADO', 'STORM SURGE' and 'HAIL'

Data Processing

Note: Before running the code the Working directory should be set to the location of "rep-data_data_StormData.csv.bz2"

```
#Import data from current directory, removing leading and trailing spaces from all values
storm_data <- read.csv(bzfile("repdata_data_StormData.csv.bz2"), strip.white=TRUE)
```

Import the data (with caching = True to avoid unnecessary lengthy repeated imports)

```
#Make all EVTYPE values upper case
storm_data$EVTYPE <- toupper(storm_data$EVTYPE)

#Remove excess spaces
storm_data$EVTYPE <- gsub("^ *|(?<= ) | *$", "", storm_data$EVTYPE, perl=T)
```

Clean up data in EVTYPE field

Simplify crop and property costs Crop and property costs are currently each split across two fields - a number and a indicated multiplier (h = hundred, k = thousand, m = million, b = billion) Below I create new columns which hold the crop and property damage cost multiplied out by the appropriate multiplier The multiplier column has some messy and a lot of missing data. Where data is not valid I assume the multiplier is 1. I have included some calculations showing the huge number of invalid multiplier values

```
#convert multiplier column to values for both CROP and PROP
#take copy of multiplier column for conversion to value
storm_data$CROPDMGEXP_as_num <- toupper(as.character(storm_data$CROPDMGEXP))
storm_data$PROPDMGEXP_as_num <- toupper(as.character(storm_data$PROPDMGEXP))

#create dataframe of valid multipliers and their values
```

```

multiplier_value <- data.frame(code = c("H", "K", "M", "B"), value = c(100, 1000, 1000000, 1000000000))

#if multiplier is not valid then use multiplier of 1
storm_data$CROPDMGEXP_as_num [!(storm_data$CROPDMGEXP_as_num %in% multiplier_value$code)] <- 1
storm_data$PROPDMGEXP_as_num [!(storm_data$PROPDMGEXP_as_num %in% multiplier_value$code)] <- 1

#Allocate valid multipliers their true value (e.g K = 1000)
storm_data$CROPDMGEXP_as_num [(storm_data$CROPDMGEXP_as_num %in% multiplier_value$code)] <- multiplier_value$value
storm_data$PROPDMGEXP_as_num [(storm_data$PROPDMGEXP_as_num %in% multiplier_value$code)] <- multiplier_value$value

#Now that characters have been replaced by numbers make column numeric in preparation for calculation
storm_data$CROPDMGEXP_as_num <- as.numeric(storm_data$CROPDMGEXP_as_num)
storm_data$PROPDMGEXP_as_num <- as.numeric(storm_data$PROPDMGEXP_as_num)

#NOTE - The calculations below show the extent of the problem of missing or invalid multiplier values -
#Invalid CROP multiplier
invalid_crop_mult <- length(storm_data$CROPDMGEXP[storm_data$CROPDMGEXP_as_num==1])
invalid_crop_mult

## [1] 618440

#Valid CROP multiplier
valid_crop_mult <- length(storm_data$CROPDMGEXP[!storm_data$CROPDMGEXP_as_num==1])
valid_crop_mult

## [1] 283857

#Invalid PROP multiplier
invalid_prop_mult <- length(storm_data$PROPDMGEXP[storm_data$PROPDMGEXP_as_num==1])
invalid_prop_mult

## [1] 466248

#Valid PROP multiplier
valid_prop_mult <- length(storm_data$PROPDMGEXP[!storm_data$PROPDMGEXP_as_num==1])
valid_prop_mult

## [1] 436049

#Percentage of CROP multipliers that are valid
valid_crop_mult / (valid_crop_mult + invalid_crop_mult) *100

## [1] 31.46

#Percentage of CROP multipliers that are valid
valid_prop_mult / (valid_prop_mult + invalid_prop_mult) *100

## [1] 48.33

```

```

#Multiply value * multiplier into new column
storm_data$CROPDMG_cost <- storm_data$CROPDMG * storm_data$CROPDMGEXP_as_num
storm_data$PROPDGMG_cost <- storm_data$PROPDGMG * storm_data$PROPDGMGEXP_as_num

event_type_summary <- aggregate(storm_data[,c("FATALITIES", "INJURIES", "CROPDMG_cost", "PROPDGMG_cost")],
#Name first column appropriately
names(event_type_summary)[1]<-"EVTYPE"

#Create total columns
event_type_summary$total_human_cost <- event_type_summary$FATALITIES + event_type_summary$INJURIES
event_type_summary$total_cost <- event_type_summary$CROPDMG_cost + event_type_summary$PROPDGMG_cost

```

Create summary table of event type (EVTYPE variable), summing the total fatalities, injuries, property cost and crop cost

Results

Across the United States, which types of events (as indicated in the EVTYPE variable) are most harmful with respect to population health? This question is easily answered using the “event_type_summary” dataframe. Although EVTYPE data is messy and inconsistently entered there are a number of clear major causes of death and injury. As can be seen in the list of top twenty causes of death and injury below ‘TORNADOS’ are the clear major cause, with ‘EXCESSIVE HEAT’, ‘TSTM WIND’, ‘FLOOD’ and ‘LIGHTNING’ having large numbers too. However you can also see the data is messy, with many events being roughly or exactly the same (e.g. ‘TSTM Wind’, ‘THUNDERSTORM WIND’, ‘HIGH WIND’, ‘THUNDERSTORM WINDS’).

```

#Return events with highest total human cost (fatalities + injuries)
ordered_human_cost <- event_type_summary[c("EVTYPE", "FATALITIES", "INJURIES", "total_human_cost")] [or

#Reorder factors before plotting
ordered_human_cost$EVTYPE <- factor(ordered_human_cost$EVTYPE, as.character(ordered_human_cost$EVTYPE))

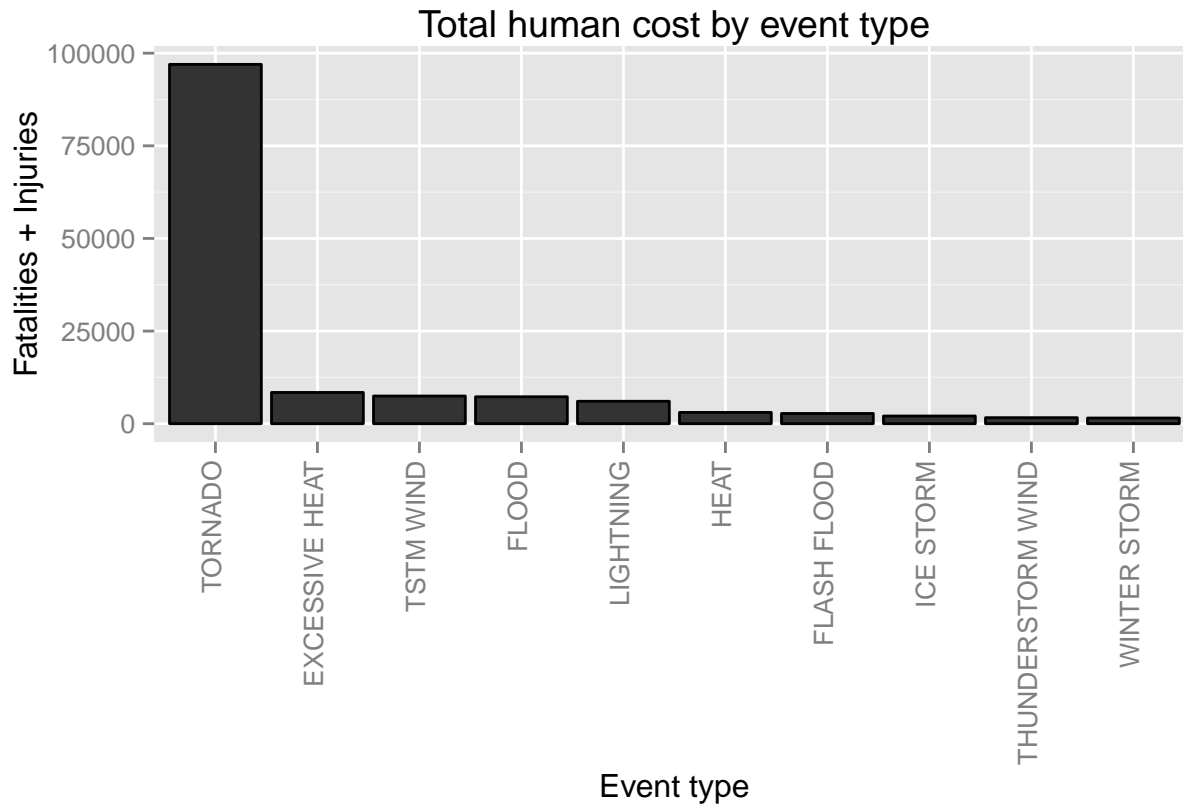
#Print the top 20 events that cause death and injury
ordered_human_cost[1:20,]

```

##	EVTYPE	FATALITIES	INJURIES	total_human_cost
## 745	TORNADO	5633	91346	96979
## 107	EXCESSIVE HEAT	1903	6525	8428
## 766	TSTM WIND	504	6957	7461
## 145	FLOOD	470	6789	7259
## 407	LIGHTNING	816	5230	6046
## 234	HEAT	937	2100	3037
## 129	FLASH FLOOD	978	1777	2755
## 376	ICE STORM	89	1975	2064
## 672	THUNDERSTORM WIND	133	1488	1621
## 873	WINTER STORM	206	1321	1527
## 309	HIGH WIND	248	1137	1385
## 203	HAIL	15	1361	1376

```
## 361 HURRICANE/TYPHOON      64    1275      1339
## 265   HEAVY SNOW          127    1021      1148
## 860   WILDFIRE            75     911       986
## 698 THUNDERSTORM WINDS    64     918       982
## 20    BLIZZARD           101     805       906
## 162    FOG                62     734       796
## 512   RIP CURRENT        368     232       600
## 858   WILD/FOREST FIRE    12     545       557
```

```
#and plot the top ten
ordered_human_cost_top <- ordered_human_cost[1:10,]
library(ggplot2)
ggplot(data=ordered_human_cost_top, aes(x=EVTYPE, y=total_human_cost)) +
  geom_bar(colour="black", stat="identity") +
  xlab ("Event type") + ylab ("Fatalities + Injuries") +
  theme(axis.text.x=element_text(angle = 90, hjust = 1, vjust = 0.3)) +
  ggtitle ("Total human cost by event type")
```



To clean up the EVTYPE data a bit I combined a number of categories if they contained certain keywords (e.g. combining all EVTYPEs that contain the text ‘TORNADO’). Note that the order of search and replace matters as there is some overlap between categories (e.g. ‘THUNDERSTORM WINDS LIGHTNING’). I have ordered the keywords from most specific to least specific to split these overlaps in the most appropriate way.

```

#Create copy of EVTYPE field
event_type_summary$EVTYPE_clean <- as.character(event_type_summary$EVTYPE)

keyword_replace <- data.frame (keyword = c("TORNADO", "HURRICANE", "TYPHOON", "LIGHTNING", "FLOOD", "HA

#Show keyword/replace table. An EVTYPE which contins the keyword anywhere within it will be replaed with
keyword_replace

```

```

##      keyword      replace
## 1    TORNADO    TORNADO
## 2   HURRICANE   HURRICANE
## 3    TYPHOON   HURRICANE
## 4   LIGHTNING   LIGHTNING
## 5     FLOOD     FLOOD
## 6      HAIL     HAIL
## 7   ICE STORM     HAIL
## 8      FOG      FOG
## 9    BLIZZARD   BLIZZARD
## 10 WINTER STORM   BLIZZARD
## 11  RIP CURRENT  RIP CURRENT
## 12     FIRE     FIRE
## 13   THUNDER THUNDERSTORM
## 14     TSTM THUNDERSTORM
## 15     WIND     WIND
## 16     HEAT     HEAT

```

```

#Search for keywords and replace value in EVTYPE_clean column
for (keyword in keyword_replace$keyword)
{
  event_type_summary$EVTYPE_clean[grepl(keyword, event_type_summary$EVTYPE_clean)] <- as.character(keyword)
}

#convert EVTYPE_clean back to factor
event_type_summary$EVTYPE_clean <- as.factor(event_type_summary$EVTYPE_clean)

```

Then we create a new EVTYPE summary table (including damage costs for analysis in second question) and rerun the human cost analysis

```

#Create event type summary based on cleaned EVTYPE data
event_type_summary_clean <- aggregate(event_type_summary[c("FATALITIES", "INJURIES", "total_human_cost")

#Name EVTYPE column
names(event_type_summary_clean)[1]<-"EVTYPE"

#Return events with highest total human cost (fatalities + injuries)
ordered_human_cost <- event_type_summary_clean[c("EVTYPE", "FATALITIES", "INJURIES", "total_human_cost")

#Print the top 20 events that cause death and injury
ordered_human_cost[1:20,]

```

```

##      EVTYPE FATALITIES INJURIES total_human_cost
## 388   TORNADO      5661    91407          97068

```

## 114	HEAT	3138	9224	12362
## 387	THUNDERSTORM	726	9448	10174
## 86	FLOOD	1525	8604	10129
## 199	LIGHTNING	817	5232	6049
## 111	HAIL	109	3457	3566
## 450	WIND	694	1979	2673
## 16	BLIZZARD	318	2159	2477
## 82	FIRE	90	1608	1698
## 165	HURRICANE	135	1333	1468
## 87	FOG	81	1077	1158
## 130	HEAVY SNOW	127	1021	1148
## 268	RIP CURRENT	577	529	1106
## 62	DUST STORM	22	440	462
## 452	WINTER WEATHER	33	398	431
## 393	TROPICAL STORM	58	340	398
## 11	AVALANCHE	224	170	394
## 78	EXTREME COLD	162	231	393
## 119	HEAVY RAIN	98	251	349
## 152	HIGH SURF	104	156	260


```

#and plot the top ten
#Reorder factors before plotting
ordered_human_cost$EVTYPE <- factor(ordered_human_cost$EVTYPE, as.character(ordered_human_cost$EVTYPE))

#Get top ten events
ordered_human_cost_top <- ordered_human_cost[1:10,]

ggplot(data=ordered_human_cost_top, aes(x=EVTYPE, y=total_human_cost)) +
  geom_bar(colour="black", stat="identity") +
  xlab ("Event type") + ylab ("Fatalities + Injuries") +
  theme(axis.text.x=element_text(angle = 90, hjust = 1, vjust = 0.3)) +
  ggtitle ("Total human cost by event type (with cleaned event types)")

```



As you can see in our cleaned dataset ‘TORNADO’ are still by far the leading cause of death and injury, followed by ‘HEAT’ with less than 1/8th the fatalities and injuries, then ‘THUNDERSTORM’ and ‘FLOOD’

Across the United States, which types of events have the greatest economic consequences?

Most of the work to answer this question has already been done. The total property damage cost and crop damage cost have been calculated, the EVTYPE variable has been cleaned up and a summary table has been created. Now it’s just a matter of seeing what the table says. NOTE: as noted above when multiplying out the damage cost, many multiplier values are missing and so the data here is very incomplete. There’s not much we can do about it on a simple pass though - perhaps the notes columns have useful information, but that’s messy and beyond the scope of this work

```
#Return events with highest total damage cost (property damage cost + crop damage cost)
ordered_cost <- event_type_summary_clean[c("EVTYPE", "CROPDMG_cost", "PROPDGMG_cost", "total_cost")] [or

#Create copy of total cost column which gives value in Billions
ordered_cost$total_cost_billions <- ordered_cost$total_cost/1e+09

#Print the top 20 events causing crop and property damage
ordered_cost[1:20,]
```

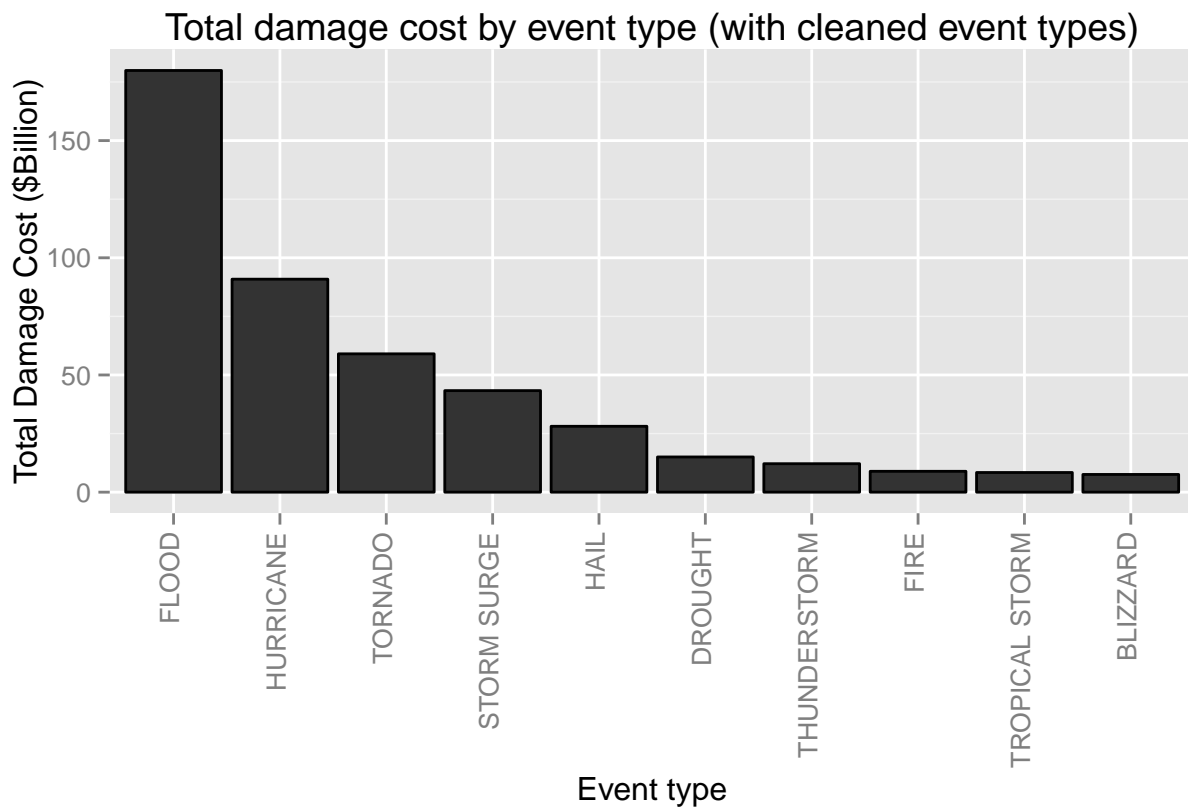
##	EVTYPE	CROPDMG_cost	PROPDGMG_cost	total_cost
## 86	FLOOD	1.238e+10	1.675e+11	1.799e+11
## 165	HURRICANE	5.516e+09	8.536e+10	9.087e+10
## 388	TORNADO	4.175e+08	5.859e+10	5.901e+10
## 318	STORM SURGE	5.000e+03	4.332e+10	4.332e+10

```
## 111          HAIL      8.134e+09    1.997e+10    2.810e+10
## 44          DROUGHT    1.397e+10    1.046e+09    1.502e+10
## 387        THUNDERSTORM 1.207e+09    1.093e+10    1.214e+10
## 82          FIRE      4.033e+08    8.497e+09    8.900e+09
## 393        TROPICAL STORM 6.783e+08    7.704e+09    8.382e+09
## 16          BLIZZARD   1.445e+08    7.414e+09    7.559e+09
## 450         WIND       7.755e+08    6.208e+09    6.983e+09
## 319        STORM SURGE/TIDE 8.500e+05    4.641e+09    4.642e+09
## 122 HEAVY RAIN/SEVERE WEATHER 0.000e+00    2.500e+09    2.500e+09
## 119         HEAVY RAIN   7.334e+08    6.942e+08    1.428e+09
## 78         EXTREME COLD  1.313e+09    6.774e+07    1.381e+09
## 100        FROST/FREEZE  1.094e+09    1.048e+07    1.105e+09
## 130        HEAVY SNOW   1.347e+08    9.326e+08    1.067e+09
## 199        LIGHTNING   1.210e+07    9.390e+08    9.511e+08
## 114         HEAT       9.045e+08    2.033e+07    9.248e+08
## 88         FREEZE      4.567e+08    2.050e+05    4.569e+08
##      total_cost_billions
## 86          179.9099
## 165          90.8725
## 388          59.0106
## 318          43.3235
## 111          28.0998
## 44          15.0187
## 387          12.1374
## 82           8.8999
## 393           8.3822
## 16           7.5589
## 450           6.9832
## 319           4.6420
## 122           2.5000
## 119           1.4276
## 78           1.3807
## 100           1.1047
## 130           1.0672
## 199           0.9511
## 114           0.9248
## 88           0.4569
```

```
#and plot the top ten
#Reorder factors before plotting
ordered_cost$EVTYPE <- factor(ordered_cost$EVTYPE, as.character(ordered_cost$EVTYPE))

#Get top ten events
ordered_cost_top <- ordered_cost[1:10,]

ggplot(data=ordered_cost_top, aes(x=EVTYPE, y=total_cost_billions)) +
  geom_bar(colour="black", stat="identity") +
  xlab ("Event type") + ylab ("Total Damage Cost ($Billion)") +
  theme(axis.text.x=element_text(angle = 90, hjust = 1, vjust = 0.3)) +
  ggtitle ("Total damage cost by event type (with cleaned event types)")
```

As you can see 'FLOOD' creates by far the largest total damage bill, followed by 'HURRICANE' with about half as much, then 'TORNADO', 'STORM SURGE' and 'HAIL'